

The Specter (and Spectra) of MEV

Guillermo Angeris*
gangeris@baincapital.com

Tarun Chitra
tarun@gauntlet.network

Theo Diamandis
tdiamand@mit.edu

Kshitij Kulkarni
ksk@eecs.berkeley.edu

August 2023

Abstract

Miner extractable value (MEV) refers to any excess value that a transaction validator can realize by manipulating the ordering of transactions. In this work, we introduce a simple theoretical definition of the ‘cost of MEV’, prove some basic properties, and show that the definition is useful via a number of examples. We also show that this definition is related to the ‘smoothness’ of a function over the symmetric group in a variety of settings and that, from this definition and some basic observations, we recover a number of results from the literature.

Introduction

Decentralized systems, such as public blockchains, allow users to submit transactions: examples of these transactions include trades, liquidations, or swaps, among any number of actions that can be performed on a blockchain. These submitted transactions are then aggregated into ordered lists, called *blocks*, by agents called validators. These validators, historically called *miners*, propose blocks to be included in the blockchain and generally have the freedom to arrange transactions within a block as they wish. There is a large amount of value that can be extracted by these privileged agents by rearranging transactions. (For example, by placing their order to purchase some asset before others are able to do the same; this idea is known as ‘front-running’.) The type of value that can be extracted by a validator in this way is termed *miner extractable value* or *MEV*, for short [DGK⁺20]. One can view the extracted value as excess returns earned by sophisticated users within these systems at the expense of unsophisticated users, which amounts to billions of dollars of excess revenue generated from MEV [Fla23]. We explore this process in more detail below.

*The authors are listed in alphabetical order.

Block construction. To understand the complexities of constructing a measure for economic fairness, we start by understanding how blocks are produced. In practice, there are multiple agents involved in the creation of a block. Users send transactions, such as trades from one asset to another on a certain exchange, via a peer-to-peer network. These transactions are received by validators—network participants who maintain the validity of state in a public blockchain. Validators put up a resource (such as energy in the case of proof of work protocols or a digital asset for proof of stake protocols) and earn fees for producing valid blocks. In most consensus protocols, validators have a temporary monopoly over both what transactions are included in a block and the order in which those transactions are executed. With this additional power, validators can increase their revenue by, for example, front-running trades, which negatively impacts users by causing them to pay higher prices than they otherwise would.

Miner extractable value. Research into MEV has generally focused on empirical properties of MEV usually found in practice. The study of MEV, and its general high-level definition, began empirically with [DGK⁺20], which identified ways in which validators could earn a profit by reordering transactions, such as trades, in ways beneficial to them. Subsequent work tended to focus on the analysis of profitability of MEV for particular types of applications [ZQT⁺21, QZG22]. These analyses focused on identifying particular types of MEV (which are identified via heuristics) and estimated the total value of this ‘value extraction’.

Formalizations. There has been some work formalizing the amount of MEV that can be extracted from users by reordering transactions in specific settings or for specific protocols [KDC22, BZ23]. These works attempt to formulate optimization problems related to measuring MEV and prove bounds on the values returned as solutions, but only apply to automated market makers such as constant function market makers [AAE⁺22]. Recent work has also demonstrated challenges with designing incentive-compatible mechanisms for allocating space on blockchains in the presence of monopolistic validators [BGR23]. Collectively, these phenomena indicate that there is a need to concretely define how the ability to order transactions enables validators to extract excess profits.

On the other hand, other forms of extractable value have not been formally analyzed. For instance, front-running in non-fungible token (NFT) auctions or liquidations in systems with leverage generate substantial MEV in practice. The main issue with such analysis is that the payoff functions for these types of MEV are non-smooth and are combinatorially complex. For instance, a liquidation can be viewed as a payoff that is one on some set of permutations of a fixed set of transactions and zero on others. Analyzing the precise expected liquidation profit relies on being able to count the number of permutations that yield a positive payoff.

Fairness. Quantifying and analyzing the impact of MEV on different classes of users has been an important research area within decentralized finance (DeFi) and the cryptocurrency space at large. Natural extensions of this work propose mechanisms for reducing or mitigating these losses. Numerous proposals have been made to mitigate MEV in both

the general consensus setting [KDL⁺21, MS22] and in specific applications [MDFO22, GSBV, CK22, BO22, JDE⁺23, RK23, Res23, Mac23, XFP23], which focus on having consensus enforce extra ordering constraints such as those that respect first-come-first-serve dynamics. This work, in comparison to the previous, mostly focuses on giving a clean theoretical definition that can be used to study MEV directly without requiring direct system constructions or mechanisms to be analyzed. Though the definition suggests some requirements for mechanisms and/or applications to have ‘low MEV’, and we note some of these requirements in this paper, we do not focus on this line of research, leaving it for future work.

This paper. In this paper, we formally define the ‘cost of MEV’ which quantifies the excess amount that a profit maximizing miner can earn from reordering a set of transactions, relative to a random ordering. The quantity essentially compares a system that is ‘order agnostic’ to one that explicitly enshrines the right to reorder transaction to validators, who seek to maximize their payoff. This definition of the cost of MEV allows us to define a notion of ‘perfectly fair functions’ that minimize this cost (these turn out to be the set of symmetric functions; *i.e.*, those invariant under permutations) and the ‘worst-cost functions’, which are essentially indicator functions of certain permutations. Next, we show that the cost of MEV shares certain smoothness properties of the underlying value function. Finally, we connect the notion of smoothness of a value function to the smoothness of functions defined on graphs, and in particular, their spectral properties. Our results allow protocol designers and theoreticians to reason about the cost of MEV in purely abstract terms. Along the way, we derive (and, in some cases, strengthen) some results already known in the literature as special cases of the provided definitions and their implications.

1 The cost of MEV

In this section, we define the cost of MEV (miner extractable value) and show some basic consequences of the definition, including simple bounds on the cost and constructing the family of functions which saturate such bounds. We then define a number of special classes of payoffs which satisfy certain properties with respect to the cost and show basic generalizations of the cost of MEV that are reasonable.

Set up. In our set up, there is an action space, defined as a set A , and a value function $f : A^n \rightarrow \mathbf{R}$, where n is the total number of transactions in the block. The set A may contain actions available to a user, including making a trade in a decentralized exchange or liquidating an underwater loan in an on-chain lending protocol. The function f characterizes the net utility accrued to a transaction validator from a particular ordered list of transactions. Some examples we will consider later are the payoff of a sandwich attack on the traders of a constant function market maker and simple liquidations. Note that we make almost no assumptions about f ; indeed, the value function f may be very complicated or difficult to write down in some practical cases.

Cost of MEV. We define the *cost of MEV* for a value function f and for some set of transactions $x \in A^n$, as

$$C(f, x) = \max_{\pi \in S} f(\pi(x)) - \mathbf{E}_{\pi \sim S} f(\pi(x)). \quad (1)$$

Here, S is the set of all possible permutations of length n (also known as the *symmetric group*) and $\pi \sim S$ means that π is uniformly randomly drawn from S . The function $C(f, x)$ characterizes how much the value of the set of transactions differs between the ‘worst case’ ordering (from the users’ perspective) and the ‘average case’ ordering. If this difference is large, then a validator may be able to reorder the transactions to extract a large amount of value from them, relative to an average ordering. Alternatively, we may view $C(f, x)$ as, roughly speaking, how much an agent with payoff f would be willing to pay in order to achieve their desired ordering, versus simply posting the transaction (assuming transactions are approximately uniformly randomly ordered in arrival times). Note that even computing this quantity might be difficult since the size of the possible permutations, $|S|$, is $n!$, which is large even for moderately small n , but we suspect many basic heuristics work very well in practice and assume this maximum can (at least approximately) be achieved.

For convenience, we will write \mathbf{E}_π for $\mathbf{E}_{\pi \sim S}$ and \max_π for $\max_{\pi \in S}$ to avoid overly burdensome notation, when the meaning is clear or easily implied.

Homogeneity. Note that the cost of MEV is homogeneous in the value function in that, for any $\alpha \geq 0$, we have

$$C(\alpha f, x) = \alpha C(f, x). \quad (2)$$

Because of homogeneity, it often suffices to pick functions f that are ‘bounded’ or normalized in some sense, for fixed x (e.g., $\max_\pi |f(\pi(x))| = 1$), since we can always rescale the result as needed.

Translation invariance. The function f and any ‘translation’ of this function, $\tilde{f}(x) = f(x) + \alpha$ both have the same cost of MEV,

$$C(\tilde{f}, x) = C(f, x),$$

for any $\alpha \in \mathbf{R}$.

Normalization. Translation invariance, together with homogeneity, means that it usually suffices to consider *normalized* functions that satisfy $0 \leq f \leq 1$ with $\max_\pi f(\pi(x)) = 1$, to derive most general properties of the cost of MEV. Since any function can be scaled and translated to satisfy these bounds, we may set

$$\tilde{f}(x) = \frac{f(x) - \min_\pi f(\pi(x))}{\max_\pi f(\pi(x)) - \min_\pi f(\pi(x))},$$

if $f(\pi(x))$ is not constant over π , for given x , and otherwise $\tilde{f}(x) = 1$. Clearly $0 \leq \tilde{f}(x) \leq 1$ and we have

$$C(f, x) = (\max_\pi f(\pi(x)) - \min_\pi f(\pi(x)))C(\tilde{f}, x).$$

Permutation invariance. We also have the property that, for any permutation $\pi \in S$, the cost of MEV is unchanged; *i.e.*,

$$C(f \circ \pi, x) = C(f, x),$$

and

$$C(f, \pi(x)) = C(f, x). \quad (3)$$

This happens because S is a group, so $\pi S = S$ for any $\pi \in S$. In this way, it makes sense to consider the set of possible transactions A^n modulo the set of possible permutations S , denoted A^n/S as the ‘natural set’ over which the cost C is measured, as opposed to measuring directly over transactions $x \in A^n$. We do not take this approach here, but do use these properties later in this section to characterize the worst-cost functions.

1.1 Worst-cost functions

In this section we introduce the worst-cost functions: a set of payoff functions that are similar in spirit to the payoffs that liquidators achieve on many decentralized finance protocols. (In some very special cases, worst-cost functions can be written as a type of ‘liquidation-like’ payoff, as in the example provided.) We show that such functions are exactly those that saturate a global bound on the cost of MEV and that they form a basis: every payoff function f can be written as a linear combination of worst-cost functions.

1.1.1 Basic bounds

For a fixed set of transactions x , it is not hard to show that, over the set of bounded functions $0 \leq f(\pi(x)) \leq 1$ for all $\pi \in S$, where there exists a π such that $f(\pi(x)) = 1$, we have that

$$C(f, x) \leq 1 - \frac{|F(x)|}{n!}, \quad (4)$$

here $F(x)$ denotes the fixed points of x under the action of S ,

$$F(x) = \{\pi \in S \mid \pi(x) = x\}.$$

(Such fixed points are sometimes known as the *stabilizers* of x by S .) To see bound (4), note that the maximum over $\pi \in S$ of $f(\pi(x))$ is achieved at some π . But, for any $\pi' \in F(x)$ we have $\pi(\pi'(x)) = \pi(x)$, so the maximum, which is at most 1 using our normalization, is achieved at at least $|F(x)|$ permutations. This means, since $f(\pi(x)) \geq 0$ for all $\pi \in S$,

$$\mathbf{E}_\pi[f(\pi(x))] \geq \frac{|F(x)|}{n!},$$

and the bound (4) follows. Note that the set $F(x)$ is nonempty since the trivial permutation is always included, so $|F(x)| \geq 1$.

Orbits. A deep—if easy to prove—statement from group theory, sometimes called the orbit-stabilizer theorem [Cla84, §2], is that

$$\frac{|F(x)|}{n!} = \frac{1}{|S(x)|},$$

where $S(x) = \{\pi(x) \mid \pi \in S\}$ is the image of x under all possible permutations S . (This set $S(x)$ is also known as the *orbit* of x under the action of the symmetric group.) We then have the bound

$$C(f, x) \leq 1 - \frac{1}{|S(x)|}.$$

In a certain sense, we may view $S(x)$ as the size of the permutations that actually matter, since some permutations simply leave x unchanged and each is counted only once in the set $S(x)$.

Implications. Given some $x \in A^n$ with $x_i \neq x_j$ for all $i \neq j$, then any permutation $\pi \in S$, except the identity (which leaves the order unchanged) will have $\pi(x) \neq x$ so $|S(x)| = n!$ and we recover the ‘simpler’ result

$$C(f, x) \leq 1 - \frac{1}{n!}.$$

More generally, if the function f is not normalized, but is nonnegative, we have the following bound, using (4) and the homogeneity of C in its first argument (2),

$$C(f, x) \leq \max_{\pi} f(\pi(x)) \left(1 - \frac{|F(x)|}{n!}\right).$$

Basic example. We will show the cost of MEV for the following toy example. Let the action space $A = \{+1, -1, L\}$, where L is a symbol indicating ‘liquidate’, while ± 1 denote ‘trades’ made with some market that increase (or decrease) the price by one unit. The payout to the liquidator, which we assume to be the validator, is given by

$$f(x) = \begin{cases} 1 & x_1 + \dots + x_k \geq n/2 \text{ and } x_{k+1} = L \\ 0 & \text{otherwise.} \end{cases}$$

(Note that, in the first case, we implicitly assume the first k trades are ± 1 .) We may interpret this function as: the validator may insert a liquidation at any point in time; they succeed at liquidating only when the price of the asset reaches a certain threshold and their liquidation happens before any other trades lower the price of the asset below their threshold. Assuming that n is even and

$$x = (1, \dots, 1, L, -1, \dots, -1),$$

with $n/2$ entries equal to $+1$ and $n/2 - 1$ entries equal to -1 , then, clearly $f(x) = 1$, so $\max_{\pi} f(\pi(x)) = 1$. On the other hand, computing $\mathbf{E}_{\pi} f(\pi(x))$ is slightly trickier. Note that

$f(\pi(x)) = 1$ if, and only if, the permutation leaves the first $n/2$ elements equal to 1, the $(n/2 + 1)$ st element equal to L , and the remaining $(n/2 - 1)!$ elements equal to -1 , in other words, when the permutation π leaves x unchanged, $\pi(x) = x$. This happens a total number of $(n/2)!(n/2 - 1)!$ ways, which gives

$$|F(x)| = (n/2)!(n/2 - 1)!,$$

or, since $f(\pi(x))$ is exactly one when $\pi(x) = x$, *i.e.*, when $\pi \in F(x)$, and zero otherwise, we have

$$\mathbf{E}_\pi f(\pi(x)) = \frac{(n/2)!(n/2 - 1)!}{n!} = \binom{n}{n/2}^{-1} \frac{2}{n},$$

which saturates the bound (4). (As we will see next, this is one example of a ‘worst-cost function’.) Putting it all together,

$$C(f, x) = 1 - \binom{n}{n/2}^{-1} \frac{2}{n} \sim 1 - 2^{-n}, \quad (5)$$

which is very close to 1 for even moderate values of n . (The right hand side follows by using the approximation $n! \sim n^n$, and the result is off by no more than a polynomial factor of n .)

This example is easily generalized to any sequence of actions $x \in A^n$ such that exactly one of the three is true: either $x_i > 0$, $x_i < 0$, or $x_i = L$ (for at most one index i) and we allow any liquidation threshold price p (compared to just $n/2$) satisfying

$$\sum_{i|x_i>0} x_i + \max_{i|x_i<0} x_i < p \leq \sum_{i|x_i>0} x_i,$$

where we ignore the single term with index k such that $x_k = L$ in the sum. In this case, the bound is very similar to the previous:

$$C(f, x) = 1 - \frac{m!(n - m - 1)!}{n!} = 1 - \binom{n}{m}^{-1} \frac{1}{n - m},$$

where m is the number of indices i such that $x_i > 0$.

1.1.2 Worst-cost functions

A particularly interesting set of functions to analyze are those which saturate the bound (4). We discuss such functions (and some consequences) here.

Fixed transactions. If a function achieves equality for the bound (4), for fixed $x \in A^n$, we will call it a *worst-cost functions for fixed transactions* x . The maximum here is achieved exactly by functions of the form

$$f(\pi(x)) = \mathbf{1}[\pi(x) = y], \quad (6)$$

where $y \in S(x)$ is some fixed element and $\mathbf{1}[\pi(x) = y]$ is the 0-1 indicator function that is 1 if $\pi(x) = y$ and 0, otherwise. To see that any (normalized) nonzero function that meets the bound must be of this form, note that $f(\pi^*(x)) = 1$ at some $\pi^* \in S$, so

$$f(\pi^* \circ \pi'(x)) = f(\pi^*(x)) = 1,$$

for every $\pi' \in F(x)$, where the first equality follows since $F(x)$ are exactly the permutations π' for which $\pi'(x) = x$. This implies that

$$\mathbf{E}_\pi[f(\pi(x))] \geq \frac{|F(x)|}{n!},$$

with equality only when $f(\pi(x))$ is zero at all permutations for which $\pi(x) \neq \pi^*(x)$. Now, if f meets the bound (4) at equality, then

$$C(f, x) = 1 - \frac{|F(x)|}{n!},$$

but, since f is normalized then $\max_\pi f(\pi(x)) = 1$ so

$$\mathbf{E}_\pi[f(\pi(x))] = \frac{|F(x)|}{n!},$$

which, from before, implies that $f(\pi(x)) = 0$ except when $\pi(x) = \pi^*(x)$, where it is 1; *i.e.*, that f can be written as

$$f(\pi(x)) = \mathbf{1}[\pi(x) = \pi^*(x)],$$

setting $y = \pi^*(x) \in S(x)$ completes the claim.

General worst-cost functions. Showing that worst-cost functions exist for any x , rather than fixing x ahead of time, is similar, but slightly more involved than the previous case. If $0 \leq f \leq 1$, then the bound of (4) still holds, but defining the class of functions f that meet this bound for all possible transactions x is somewhat more delicate.

In this case, we define the set of equivalence classes of the action space A^n under the possible permutations, which we denote

$$A^n/S = \{S(x) \subseteq A^n \mid x \in A^n\},$$

Note that the sets A^n/S form a partition of A^n as they are pairwise disjoint and their union is all of A^n , which follows from the fact that S is a group. Additionally, since $x \in S(x)$, each set is nonempty. The set of functions that saturate the bound for all $x \in A^n$, *i.e.*, the global *worst-cost functions* are then exactly those of the form

$$f(x) = \sum_{Q \in A^n/S} \mathbf{1}[x = x_Q], \tag{7}$$

where $x_Q \in Q$ for each $Q \in A^n/S$. This (potentially uncountable) sum is justified since, for any x , there is at most one term that will have value exactly equal to one. We can view

the function f as picking one ‘canonical’ permutation for every equivalence class $Q \in A^n/S$ and assigning it a value of 1 if $x \in Q$ and matches this permutation exactly. Otherwise, it assigns a value of 0. That this family of functions is nonempty, for general sets A^n , requires the axiom of choice, but these sets are generally far more structured in practice, so it should be possible to construct this worst-cost function directly.

Proof. The proof that every normalized worst-cost function is exactly of this form follows from viewing f restricted a single equivalence class $Q \in A^n/S$. Let f be a function of the form of (7), then, for every $x \in Q$,

$$f(x) = \mathbf{1}[x = x_Q],$$

since x is in exactly one equivalence class Q . Of course, by definition of Q , we have that $\pi(x) \in S(x) = Q$ for every $\pi \in S$, so

$$f(\pi(x)) = \mathbf{1}[\pi(x) = x_Q],$$

which we know are exactly the worst-cost functions for a fixed $x \in A^n$ from the previous discussion. Since

$$C(f, x) = C(f, y)$$

for any $y \in S(x) = Q$ by using the permutation invariance of C in (3), then this is also a worst-cost function for any $y \in Q$. Finally, summing over all possible Q and noting that at most one such indicator has value 1 for each Q , and therefore for each $x \in A^n$, gives the final result.

Worst-cost functions as a basis. An interesting fact is that the set of all worst-cost functions is a basis for all functions f . In particular, we have that $f(\pi(x))$, viewed as a function of π and holding $x \in A^n$ constant, can be written as a linear combination of worst-cost functions

$$f(\pi(x)) = \sum_{y \in S(x)} f(y) \mathbf{1}[\pi(x) = y]. \quad (8)$$

This expression is well-defined since $S(\pi(x)) = S(x)$ for any permutation π , by definition. It is not hard to see that equation (8) is true. Consider $f(\pi(x))$, then $\pi(x) \in S(x)$ so exactly one term in the sum is nonzero: the one corresponding to $f(\pi(x))$. The argument is easy to generalize to the broader case where the transactions x are not fixed and in that case we have

$$f(x) = \sum_{Q \in A^n/S} \sum_{y \in Q} f(y) \mathbf{1}[x = y],$$

and we can see that this indeed corresponds to f by (carefully!) interpreting the sums.

1.2 Fair and unfair functions

In this section we define perfectly fair functions, which is a strong requirement that a few value functions meet in practice. We show that unfair functions, those which have high cost of MEV, satisfy a number of bounds of possible interest. For example, if a normalized function is very ‘localized’ (*i.e.*, it takes on very large values at only a small number of positions and small values elsewhere) then it is very unfair. We also show a partial converse: any ‘large enough’ function that is has high cost must be reasonably localized.

Perfectly fair functions. Since the maximum is always no smaller than the expectation, we have that

$$C(f, x) \geq 0.$$

We say a value function is *perfectly fair for transactions* $x \in A^n$ if

$$C(f, x) = 0.$$

Note that f is perfectly fair for transactions x if, and only if,

$$f(\pi(x)) = f(x),$$

for all $\pi \in S$. Based on this, we will say it is *perfectly fair* if

$$C_s(f) = \sup_{x \in A^n} C(f, x) = 0,$$

since $C_s(f) = 0$ implies that $C(f, x) = 0$ for all $x \in A^n$. Or, equivalently, the perfectly fair functions are the set of symmetric functions, defined as those which have $f \circ \pi = f$ for all $\pi \in S$. While this might seem like an overly-burdensome restriction, there do exist examples of perfectly fair functions in practice, these include those of concave pro-rata games [JDE⁺23] or sealed-bid auctions [Rou16], both of which are mechanisms that do not depend on the order in which user actions were received. We will see how to construct a ‘perfectly fair’ function (in expectation) from a general payoff function f in the extensions.

Unfair functions. We can push some of the techniques described here to create lower bounds for general functions f and fixed transactions x . In particular, if a bounded function $f \leq 1$ satisfies $f(\pi(x)) \geq \beta$ for some $\pi \in H \subseteq S$ and $f(\pi(x)) \leq \alpha$ for each $\pi \in S \setminus H$ with $0 \leq \alpha \leq \beta \leq 1$, then $\max_{\pi} f(\pi(x)) \geq \beta$ while

$$f(\pi(x)) \leq \mathbf{1}[\pi \in H] + \alpha \mathbf{1}[\pi \in S \setminus H],$$

for every $\pi \in S$, so

$$\mathbf{E}_{\pi}[f(\pi(x))] \leq \alpha + \frac{|H|}{n!}(1 - \alpha).$$

Putting this all together gives the following lower bound on $C(f, x)$,

$$C(f, x) \geq \beta - \alpha - \frac{|H|}{n!}(1 - \alpha). \tag{9}$$

We can think of the set $H \subseteq S$ as the set of permutations under which f is ‘spiky’ for transactions x . If this set H is (a) small relative to S and (b) has values that are bounded away from the remaining permutations with a large margin $\beta \gg \alpha$, for instance, then the cost is always guaranteed to be large.

A particularly interesting case to analyze is when $\alpha = 0$ and $f(\pi(x))$ is nonconstant over π . In this case, $f(\pi(x))$, seen as a function over $\pi \in S$, has support included in H and, since the maximum is achieved somewhere in H , we must have that $\beta = 1$, assuming f is normalized. This gives the simple bound:

$$C(f, x) \geq 1 - \frac{|H|}{n!},$$

Note that this bound is tight against (4) exactly when $H = F(x)$, the set of stabilizers of x under S , which is, when f is normalized, whenever $f(\pi(x)) = \mathbf{1}[\pi(x) = \pi'(x)]$ for some fixed permutation $\pi' \in S$, giving another proof for the form of worst-cost functions. So, if the support is smaller than the set of all permutations, $|H| < n!$, then necessarily we have positive cost of MEV. If the function f is not normalized but is nonnegative with $f(\pi(x)) \geq 0$ for all $\pi \in S$, we have that

$$C(f, x) \geq \left(1 - \frac{|H|}{n!}\right) \max_{\pi \in S} f(\pi(x)),$$

by the homogeneity of C in its first argument (2).

Partial converse. We can also provide a partial converse to the above that says: if a normalized function has high cost of MEV, then the function ‘almost looks like’ a worst-cost function. More explicitly, we have the following two claims. First, if $C(f, x) \geq \alpha$, for a fixed list of transactions x , then

$$f(\pi(x)) \geq \alpha \mathbf{1}[\pi(x) = y],$$

for some $y \in S(x)$; that is, the function f is at least as large as some indicator. This result follows pretty much directly from the definition: there is some π such that $f(\pi(x)) \geq \alpha$ and setting $y = \pi(x) \in S(x)$ suffices. Of course, the function f could potentially be no smaller than an indicator over a much larger set, call it $T \subseteq S(x)$, but the following statement bounds the size of the resulting set. If, for some $\eta > 0$ we have that

$$f(\pi(x)) \geq \eta \mathbf{1}[\pi(x) \in T],$$

then the size of the set T must be bounded by

$$|T| \leq \frac{(1 - \alpha)n!}{\eta |F(x)|}.$$

That is, fixing $\eta > 0$: if the cost, bounded by α , is very large ($\alpha \approx 1$) then the set T must be very small for reasonable values of η . Another way of stating this is: if the function has

high cost, then its largest values are very localized. (Where the most localized functions with $|T| = 1$ are exactly the worst-cost functions.) To see this, if f is normalized and $f(\pi(x)) \geq \eta \mathbf{1}[\pi(x) \in T]$, we have

$$\alpha \leq C(f, x) \leq 1 - \frac{1}{n!} \eta |F(x)| |T|,$$

and the result follows by rearranging.

1.3 Discussion and generalizations

From the previous discussion on the cost of MEV, we may view the cost of MEV, C , as measure of how ‘spiky’ or ‘localized’ a particular payoff function f is, with respect to the permutations; the higher the cost, the more localized such a function is, and vice versa. (We will see another connection to this interpretation, via the Fourier transform on graphs, in a later section.) We may also view the cost of MEV as: how much is a user willing to pay (in expectation) in order to ensure that their preferred ordering is achieved on chain? Alternatively, we may view the function C as a type of ‘efficiency’ metric for MEV, when the functions are normalized: values very close to 1 denote a near-maximal amount of miner extractable value possible and values very close to 0 denote small amounts. Of course, there are many other reasonable measures of the cost of MEV that are very similar in spirit to the one presented above. We outline some here.

Ratio. Another reasonable metric to measure the cost of MEV is a ‘ratio’ as opposed to an absolute difference:

$$\tilde{C}(f, x) = \frac{\mathbf{E}_\pi f(\pi(x))}{\max_\pi f(\pi(x))}.$$

For normalized functions this reduces to $\mathbf{E}_\pi f(\pi(x))$ and is homogeneous of degree zero in its function argument. Note also that it is at most 1 for any function f , and is bounded between 0 and 1 for nonnegative f . Many results have direct analogues in this case, but we find that working with the additive form is simpler and has a number of other interesting results. For example, this multiplicative variant, \tilde{C} , is not translation invariant, nor does it satisfy the smoothness results presented in the next section.

Randomness. Many papers have pointed out the importance of randomness in achieving potentially fair mechanisms. This requires extending the current definition slightly to functions $f : A^n \times \Omega \rightarrow \mathbf{R}$ where \mathbf{P} is a probability distribution over the sample space Ω , in the following way:

$$\bar{C}(f, x) = \max_{\pi \in P} \mathbf{E}_{\omega \sim \mathbf{P}} f(\pi(x), \omega) - \mathbf{E}_\pi [\mathbf{E}_{\omega \sim \mathbf{P}} f(\pi(x), \omega)].$$

The order of the operations has the following interpretation: the mechanism uses some randomness, unknown to the user until the ordering π has been provided, and the user wishes to compare the expected value of an ordering that maximizes their cost function

when compared to a uniformly chosen ordering. This is just a special case of (1) where the function being evaluated is $\bar{f} : A^n \rightarrow \mathbf{R}$ where \bar{f} is

$$\bar{f}(x) = \mathbf{E}_{\omega \sim \mathbf{P}} f(x, \omega),$$

though the practical interpretation differs substantially.

An aside: constructing ‘fair’ functions. With this extension, we can convert any (deterministic) function $f : A^n \rightarrow \mathbf{R}$ to a ‘perfectly fair’ function, in expectation, by letting the sample space be the set of permutations, $\Omega = S$, letting the distribution be uniform over the sample space, $\mathbf{P}(\omega) = 1/n!$, and by setting

$$\tilde{f}(x, \omega) = f(\omega(x)).$$

From before, this function has the interpretation that the protocol chooses a uniformly random permutation ω after any set of transactions x has been provided and permutes these transactions with the randomly chosen permutation to get $\omega(x)$. To prove fairness, note that, for any $\pi \in P$,

$$\mathbf{E}_{\omega \sim \Omega} f(\omega(\pi(x))) = \mathbf{E}_{\omega \sim \Omega} f(\omega(x)),$$

since $\omega \circ \pi$ is also uniform over S , so

$$\bar{C}(f, x) = 0.$$

and therefore

$$\bar{C}_s(f) = \sup_{x \in A^n} \bar{C}(f, x) = 0.$$

for any $x \in A^n$.

Local vs. global bounds. In a certain sense, lower bounds on C_s are easy: show that there exist some set of transactions for which the cost C is large. We may then view lower bounds as ‘local’ in that a single example suffices to show that a lower bound on the cost holds. Upper bounds are, on the other hand, much harder, as we must make a claim about all possible lists of transactions from the set of actions A^n , which we may view as a type of global statement.

Going forward. From here, it is not clear that a lot more can be said in a general setting without one of two assumptions: either (a) assume some structure about the set A (such as, *e.g.*, metric structure, or some notion of smoothness, which will play in various ways with the symmetric group), or (b) assume that the transactions $x \in A^n$ are fixed and, instead, look at the properties of C from a ‘local’ perspective. This latter view has deep connections to spectral graph theory and Fourier transforms over groups. In a sense, we can view the former statement as being what we can say over all transactions, whereas the latter property is what we can say over all permutations. We describe each approach in the following two sections.

2 Smoothness and permutations

We now introduce the ‘global’ approach to analyzing the cost of MEV of a value function over all transactions in A^n . To do this, we use smoothness properties of the value function and some metric structure on the set of transactions. We will assume that there is a subset $B \subseteq A^n$ closed under permutations, *i.e.*, $\pi B = B$ for all $\pi \in S$ and that $f : B \rightarrow \mathbf{R}$. If B can be endowed with a metric $d : B \times B \rightarrow \mathbf{R}_+$ that is *permutation independent*, that is, if, for any $\pi \in S$, we have

$$d(\pi(x), \pi(y)) = d(x, y),$$

then there are useful smoothness results that hold for C . (Examples of permutation-independent metrics include those induced by the standard norms, for example.)

2.1 Smoothness bounds

In general, C has similar smoothness characteristics as the function f over B . For example, if f is L -Lipschitz-continuous over B , that is, if for $L \geq 0$, we have

$$|f(x) - f(y)| \leq Ld(x, y),$$

for any $x, y \in B$, then $C(f, \cdot)$ is $2L$ -Lipschitz-continuous in that,

$$|C(f, y) - C(f, x)| \leq 2Ld(x, y). \quad (10)$$

Proof. This is easy to see since, for the expectation, we have

$$|\mathbf{E}_{\pi'} f(\pi'(x)) - \mathbf{E}_{\pi} f(\pi(y))| \leq \mathbf{E}_{\pi} |f(\pi(x)) - f(\pi(y))| \leq \mathbf{E}_{\pi} [Ld(\pi(x), \pi(y))] = Ld(x, y),$$

while for max we have

$$f(\pi'(x)) - \max_{\pi} f(\pi(y)) \leq \max_{\pi} |f(\pi(x)) - f(\pi(y))| \leq L \max_{\pi} d(\pi(x), \pi(y)) = Ld(x, y)$$

for all $\pi' \in S$ so

$$|\max_{\pi} f(\pi(y)) - \max_{\pi'} f(\pi'(x))| \leq Ld(x, y).$$

Applying the triangle inequality and using the definition of C yields the final result (10).

Tightness. This bound is tight in that the constant cannot be improved for general functions f , sets B , and metrics d , unless other assumptions are made. We can see this from the simple example where $B = A^n = \mathbf{R}^n$, the metric d is the one implied by the ℓ_1 norm, and the value function is

$$f(x) = x_1 - x_2 - \cdots - x_n.$$

First, we see that f is 1-Lipschitz continuous in the ℓ_1 norm, $\|\cdot\|_1$, that is:

$$|f(x) - f(y)| \leq \|x - y\|_1 = \sum_{i=1}^n |x_i - y_i|,$$

since

$$|f(x) - f(y)| = |(x_1 - y_1) - (x_2 - y_2) - \cdots - (x_n - y_n)| \leq \|x - y\|_1,$$

by applying the triangle inequality n times. Now, consider $x = 0$ and $y = e_1 = (1, 0, \dots, 0)$. Since, $\|x - y\|_1 = \|y\|_1 = 1$ the smoothness bound (10) implies that

$$|C(f, x) - C(f, y)| \leq 2\|x - y\|_1 = 2, \quad (11)$$

where we have used the fact that f is 1-Lipschitz continuous with respect to the ℓ_1 -norm. We will show that this bound on the cost (11), implied by the smoothness bound (10) is saturated in a certain limit. First, consider $f(\pi(x))$ where $x = 0$, where we have

$$f(\pi(x)) = 0$$

for any $\pi \in S$ so $C(f, x) = 0$. Now, consider $y = e_1$. Here, where we have $\max_{\pi} f(\pi(y)) = 1$ and

$$\mathbf{E}_{\pi} f(\pi(y)) = \frac{1}{n} - \left(1 - \frac{1}{n}\right) = \frac{2}{n} - 1.$$

We can see this since $f(\pi(y)) = 1$ only when π fixes the first element of y to be 1, which happens with probability $1/n$, and takes the value -1 at all other times, with probability $1 - 1/n$. Using the definition of $C(f, y)$ we then get

$$C(f, y) = 2 \left(1 - \frac{1}{n}\right),$$

so

$$|C(f, x) - C(f, y)| = |C(f, y)| = 2 \left(1 - \frac{1}{n}\right).$$

(Compare this with the bound (11).) Taking $n \uparrow \infty$, we see that the cost of MEV approaches 2, which saturates the bound provided by (10) in the limit. (Scaling this choice of f by any value $L > 0$ shows that this bound is saturated for any L .)

2.1.1 Global upper bounds

If we know the function is smooth and the set of transactions B is bounded with diameter at most t , *i.e.*,

$$t \geq \sup_{x, y \in B} d(x, y),$$

then we have the immediate upper bound on the maximum cost

$$C_s(f) = \sup_{x \in B} C(f, x),$$

given by:

$$C_s(f) \leq 2Lt + \inf_{x \in B} C(f, x). \quad (12)$$

So, from the smoothness of f we receive a general bound on the cost of MEV that matches our intuition: if the function f is ‘smooth’ over the transactions, even independent of possible ordering effects, then it has ‘low’ cost. Additionally, if any action $x \in B$ has $x_i = x_j$ for every $i, j = 1, \dots, n$, then the second term in the right-hand-side of (12) is zero, so

$$C_s(f) \leq 2Lt.$$

The bound (12) generalizes and improves upon some known results over constant function market makers [CAE22, KDC22], and we show some explicit examples below.

2.1.2 Extensions

Note that we do not use the fact that d is a metric in any of the above derivations, so it may be any function which satisfies the Lipschitz-like condition

$$|f(x) - f(y)| \leq Ld(x, y).$$

Additionally if d is not permutation independent, we may define

$$\tilde{d}(x, y) = \max_{\pi} d(\pi(x), \pi(y)),$$

which is permutation independent and satisfies $\tilde{d} \geq d$. All bounds above hold over this new ‘distance-like’ function, \tilde{d} .

2.2 Smooth mechanism examples

We show some basic applications of the smoothness upper bound (12) to frontrunning and sandwiching in decentralized exchanges.

2.2.1 Frontrunning

A classic example of MEV is the notion of *frontrunning*; *i.e.*, the fact that parties (such as validators) are willing to pay some cost in order to be the first transaction executed.

Action space. In this particular case, we will assume that the set of possible actions is

$$A = \underbrace{(\mathbf{R}_+ \times \{0\})}_{\text{actions by traders}} \cup \underbrace{(\{0\} \times \{\delta\})}_{\text{actions by validator}},$$

where $\delta \in \mathbf{R}_+$ is some nonnegative number. We can view a tuple $(x, y) \in A$ as an action, where x (a trade made by a trader) is nonzero only when y (a trade made by the validator) is zero, and $y = \delta$ is the only possible trade that can be made by a validator. We then constrain the possible actions to be bounded in the following sense:

$$B = \{((x_1, y_1), \dots, (x_n, y_n)) \in A^n \mid \|x\|_1 \leq M \text{ and exactly one entry of } y \text{ is nonzero}\}.$$

In other words, the total volume traded is constrained to be at most M , while the validator only submits exactly one trade of size δ . To avoid overly burdensome notation, we will overload symbols slightly and write $(\tilde{x}, \tilde{y}) \in B$ to mean

$$((\tilde{x}_1, \tilde{y}_1), \dots, (\tilde{x}_n, \tilde{y}_n)) \in B.$$

In other words, \tilde{x} is the vector of trades by traders and \tilde{y} is the vector of trades by the validator. By definition, \tilde{x}_i is nonzero only when $\tilde{y}_i = 0$ and vice versa.

Metric. In this case, we will simply set the ‘metric’ d to be given by the following function, defined for some

$$z = (\tilde{x}, \tilde{y}) \in B,$$

as

$$d(z, z') = \sum_{i=1}^n \max\{|\tilde{x}_i|, |\tilde{x}'_i|, |\tilde{x}_i - \tilde{x}'_i|\},$$

where \tilde{x}' has the same meaning as \tilde{x} , except for action z' . In particular, this ‘metric’ is permutation-independent and is not a true metric. It measures something very similar, but not quite equivalent, to the total volume of trades on the CFMM made by traders.

Diameter bound. Note that B , the set of allowable actions, has ‘diameter’ bounded with respect to this ‘metric’ using the definition of the set B and so satisfies

$$d(z, z') \leq 2M \tag{13}$$

for every $z, z' \in B$. To see this, apply the triangle inequality,

$$|\tilde{x}_i - \tilde{x}'_i| \leq |\tilde{x}_i| + |\tilde{x}'_i|,$$

which implies that

$$d(z, z') \leq \sum_{i=1}^n \max\{|\tilde{x}_i|, |\tilde{x}'_i|, |\tilde{x}_i| + |\tilde{x}'_i|\} = \sum_{i=1}^n |\tilde{x}_i| + |\tilde{x}'_i| \leq 2M.$$

Surprisingly, this bound is tight when the number of trades is greater than one. We can see this by setting $\tilde{x}_i = 2M/n$ for even indices i and 0 otherwise, and $\tilde{x}'_i = 2M/n$ for odd indices i and 0 otherwise. Clearly $\|\tilde{x}\|_1 \leq M$ and similarly for \tilde{x}' yet we have that $d(z, z') = 2M$. (This assumes that the number of entries, n , is even and greater than one, otherwise one may scale the nonzero entries of \tilde{x}' , since $\|\tilde{x}'\|_1 = M(1 - 2/n)$ when $n \geq 2$ is odd, to match the bound.)

Payoff function. We assume that the market acts like a two-asset fee-free constant function market maker [AAE⁺22], which we define as a nondecreasing concave function $G : \mathbf{R}_+ \rightarrow \mathbf{R}_+$ with $G(0) = 0$. (This function G is sometimes called the forward exchange function [AAE⁺22, §4.1].) We assume the market begins at some state 0 and, assuming some sequence of trades $t_1, \dots, t_k \in \mathbf{R}_+$ were made, then trade $k + 1$, with some input amount t_{k+1} , pays out

$$G(t_1 + \dots + t_k + t_{k+1}) - G(t_1 + \dots + t_k). \quad (14)$$

In this case, the payoff for the validator is as follows. Given some actions $z = (\tilde{x}, \tilde{y}) \in B$, let k be the index such that $\tilde{y}_k = \delta$, then the payoff for the validator is

$$f(z) = G(\tilde{x}_1 + \dots + \tilde{x}_{k-1} + \delta) - G(\tilde{x}_1 + \dots + \tilde{x}_{k-1}).$$

Smoothness. Since G is concave, we have that

$$G'(t) \leq G'(0),$$

for $t \geq 0$, so we have, for $t' \geq 0$,

$$|G(t) - G(t')| \leq G'(0)|t - t'|.$$

This implies that, given two actions $z, z' \in B$ and setting

$$t = \sum_{i=1}^k \tilde{x}_i, \quad t' = \sum_{i=1}^{k'} \tilde{x}'_i,$$

where k is the index for which $\tilde{y}_k = \delta$, *i.e.*, when the validator makes their trade in action z , and k' is similar for action z' , we have

$$\begin{aligned} |f(z) - f(z')| &= |G(t + \delta) - G(t) - (G(t' + \delta) - G(t'))| \\ &\leq |G(t + \delta) - G(t' + \delta)| + |G(t) - G(t')| \leq 2G'(0)|t - t'|. \end{aligned} \quad (15)$$

Using the triangle inequality, we have,

$$|t - t'| \leq \sum_{i=1}^k |\tilde{x}_i - \tilde{x}'_i| + \sum_{i=k+1}^{k'} |\tilde{x}'_i| \leq d(z, z').$$

Here we assumed $k \leq k'$ but swapping z for z' gives the complete claim. We then get the final result using (15) and the above:

$$|f(z) - f(z')| \leq 2G'(0)d(z, z').$$

Cost of MEV. Using the bound (12) and the ‘diameter’ bound (13), we then get that the cost of MEV is at most

$$C_s(f) \leq 8G'(0)M,$$

where we used the fact that $\inf_{x \in B} C(f, x) = 0$ by setting all traders’ trades to be zero. This bound matches our intuitions: as the size of others’ trades grows larger (M is large) the cost of MEV similarly becomes larger, roughly linearly in the total amount of volume traded. (Roughly speaking, we expect that the larger the total volume of others’ trades, the more a validator would be willing to pay to front-run.) This bound is only reasonable when the validator’s trade δ is at least on the order of the total trade volume M ; otherwise, the bound is very loose when δ is very small.

2.2.2 Sandwiching

Another example of MEV in decentralized exchanges is the notion of *sandwich attacks*, *i.e.*, an adversary (such as the validator introduced above) inserts a buy transaction before users purchase an asset and a sell transaction after. Note that, intuitively, the amount of profit generated is always nonnegative since the adversary always ‘purchases’ an asset at a price no higher than she sells it at.

Action space. The set of possible actions in this setting takes the form

$$A = (\mathbf{R}_+ \times \{0\}) \cup (\{0\} \times \mathbf{R})$$

Similar to the previous example, a tuple $(x, y) \in A$ is an action where x is a trade made by a trader, which is nonzero only when y (the trade made by the adversary) is zero. Alternatively, we can have either $y > 0$ or $y < 0$ as the possible choices for the sandwicher’s trade, which can be interpreted as buy and sell orders in the decentralized exchange, respectively.

Sandwiching mechanics. In our model, we will assume that the sandwicher observes the trades to be included in the block, which we call $x_1, \dots, x_k \in \mathbf{R}_+$. The sandwicher then submits a purchase order of $\delta \in \mathbf{R}_+$, ideally inserted before all other trades happen, and submits a buy order of $\gamma \in \mathbf{R}_+$, ideally inserted after all other trades happen. We assume the same model of the market as the previous example, defined by some function $G : \mathbf{R} \rightarrow \mathbf{R}$, which is strictly increasing and concave, with $G(0) = 0$. (We allow negative trades in this case to indicate an amount which must be tendered instead of received.) Assuming the sandwicher purchases some amount of an asset by tendering δ and then sells back no more than the amount $G(\delta)$ received, the trades must satisfy the inequality

$$G(t + \delta) - G(t + \delta - \gamma) \leq G(\delta) - G(0) = G(\delta), \quad (16)$$

where $t = x_1 + x_2 + \dots + x_k$ is the total of all other trades. Note that this uses the definition (14) of the market G , and assumes that the trade δ is made in the first position and γ is made in the last. (It is a useful exercise to show that, due to the concavity of G ,

and the fact that all trades are nonnegative, this is the largest possible δ and γ that can be tendered, or received, respectively.) The left hand side can be interpreted as the amount of γ received for tendering $G(\delta)$ back to the market. The inequality simply says that this amount, interpreted as the amount tendered, is bounded by the amount of the asset received by the initial trade of δ (*i.e.*, the sandwicher does not tend more to the market than the amount that was received from the initial trade of δ).

Constrained action space. We will define the set of possible actions B in the following way:

$$B = \{((x_1, y_1), \dots, (x_n, y_n)) \in A^n \mid \|x\|_1 \leq M \text{ and exactly one pair } i, j \text{ with } y_i > 0 \text{ and } y_j < 0 \text{ which must satisfy (16)}\}.$$

In other words, B is defined as the set of trades for which the total trade volume, excluding the sandwicher's trades, is no more than M , and the only two trades by the sandwicher, given by, say, $y_i > 0$ and $y_j < 0$, satisfy (16) in the following way:

$$G(t + y_i) - G(t + y_i + y_j) \leq G(y_i), \quad (17)$$

where $t = x_1 + x_2 + \dots + x_n$.

Value function. Overloading notation in the same way as the previous example, we may write an action $z = (\tilde{x}, \tilde{y}) \in B$. We define the value function, or payoff, for this action (from the perspective of the sandwicher) to be the amount received of the asset minus the amount tendered originally, which is simply

$$f(z) = -(\tilde{y}_j + \tilde{y}_i). \quad (18)$$

(The negative sign comes from the fact that negative values denote amounts received from the market.)

Metric. Finally, we set the 'metric' d for $z = (\tilde{x}, \tilde{y}) \in B$ and $z' = (\tilde{x}', \tilde{y}') \in B$, where we have overloaded notation in the same way as the previous example, to

$$d(z, z') = \max\{\|\tilde{x}\|_1, \|\tilde{x}'\|_1\}.$$

In other words, this 'metric' simply measures the largest volume of trades from the two possible actions. By definition of B we have that the 'diameter' is no more than

$$d(z, z') \leq M.$$

Smoothness. Using the definition of B , the smoothness is very easy to prove. Let $z = (\tilde{x}, \tilde{y}) \in B$ be an action. Note that, since G is increasing, then, using (17) we have

$$G(t + \tilde{y}_i + \tilde{y}_j) \geq G(t + \tilde{y}_i) - G(t) \geq 0,$$

where $t = \tilde{x}_1 + \dots + \tilde{x}_n$. This implies that

$$t + \tilde{y}_i + \tilde{y}_j \geq 0,$$

or that the payoff for any action is bounded from above by

$$f(z) = -(\tilde{y}_i + \tilde{y}_j) \leq t \leq \|\tilde{x}\|_1.$$

Note that this inequality is rather intuitive: it states that the maximum amount of MEV available from sandwiching cannot be more than the total volume of trades happening on the market. Since, from before, the payoff is always nonnegative, we have that $f(z), f(z') \geq 0$ so,

$$|f(z) - f(z')| \leq \max\{f(z), f(z')\} \leq \max\{\|\tilde{x}\|_1, \|\tilde{x}'\|_1\} = d(z, z'),$$

which means that f is 1-smooth with respect to the ‘metric’ d .

Cost of MEV. Using (12), and the fact that the ‘diameter’ of B is no more than M , we find that the worst-case cost of MEV is

$$C_s(f) \leq 2M.$$

This immediately strengthens known bounds for $C_s(f)$ from [KDC22], which had $C_s(f) = O(\log n)$ versus $C_s(f) = O(1)$, assuming the market volume is constant, as we have here. Again, this is reasonably intuitive: the maximum cost of MEV cannot be more than a constant factor away from the total market volume. In this particular case, using the fact that f is nonnegative, we may actually strengthen this result to

$$C_s(f) \leq M,$$

by directly using the definition of the cost of MEV.

3 Permutations on graphs

In this section, we will show the connections between the function $f(\pi(x))$ (where the permutation $\pi \in S$ is free while the transaction $x \in A^n$ is fixed) and a graph whose vertices are indexed by the set of permutations S , with edges encoding particular structure that we define next. We will show a number of basic results which generalize some of those known in the literature.

3.1 Definitions and examples

We will define a *permutation graph* as a connected, undirected graph $G = (V, E)$ where the vertices V are indexed by the set of permutations S and the edge set E is composed of subsets of V of cardinality 2. While the vertices of the permutation graph are fixed in this definition, the edges are not and we are free to choose E in any way we like, so long as the graph G is connected. (We will make use of this freedom later when analyzing f .) For convenience, we will overload notation slightly by defining:

$$f_i = f(\pi_i(x)), \quad i = 1, \dots, n!,$$

for fixed $x \in A^n$ and where π_i ranges over all permutations. Unless stated otherwise, we will assume that x is fixed and implicitly included in the definition of the f_i for the remainder of the paper. We will also identify $f \in \mathbf{R}^{n!}$ as an $n!$ -vector such that quantities like the ℓ_2 norm, $\|f\|_2$, have the obvious interpretations.

A basic bound. While this definition appears trivial, there are already a number of useful consequences. For example, for any path $P \subseteq E$ from node i to node j ,

$$|f_i - f_j| \leq \sum_{\{k, \ell\} \in P} |f_k - f_\ell|,$$

which follows from the triangle inequality. Taking the maximum over i and j and a shortest path for each pair, we have that

$$\max_{i, j \in V} |f_i - f_j| \leq \mathbf{diam}(G) \left(\max_{\{k, \ell\} \in E} |f_k - f_\ell| \right), \quad (19)$$

where $\mathbf{diam}(G)$ is the diameter of the graph; *i.e.*, the largest distance between any two nodes of the graph, where the distance is defined as the length $|P|$ of the shortest path P between the nodes. (This is always finite as we assumed the graph G is connected.) Since

$$C(f) \leq \max_{i, j \in V} |f_i - f_j|,$$

this also immediately implies a bound on the cost of fairness. Inequality (19) already gives us an idea for why a ‘good’ choice of E is essential: in order to get a tight bound, we would like to choose E to be ‘simple enough’ to get a reasonable bound on the maximum difference between adjacent permutations, while making it connected enough so the diameter of the graph, $\mathbf{diam}(G)$, is small. We outline a few important examples below.

Complete graph. Perhaps the simplest example of an edge set E is the set of all possible edges, in which case the graph G is the complete graph of $n!$ vertices. In this case bound (19) is trivial as $\mathbf{diam}(G) = 1$ (since any node is exactly one hop away from any other node) while

$$\max_{\{k, \ell\} \in E} |f_k - f_\ell| = \max_{k, \ell \in V} |f_k - f_\ell|,$$

which is what we wanted to bound in the first place. In almost all cases, the complete graph almost universally yields weak and/or trivial bounds for many quantities, so it is only useful as an example of the definitions.

Transposition graph. A second possible choice of edge set E is for every pair of nodes i and j , we add an edge $\{i, j\}$ to E if π_i and π_j differ by exactly one transposition. (In other words, if there are two elements in π_i which, when swapped, results in π_j .) We will call this graph the *transposition graph*. This graph is connected and has diameter $\mathbf{diam}(G) = n - 1$ since any permutation can be written as $n - 1$ transpositions (*cf.*, [Con]). If we have any bound on the maximum difference of f in transpositions, say $|f_i - f_j| \leq C$ if $\{i, j\} \in E$, then, applying (19), we get:

$$\max_{i,j \in V} |f_i - f_j| \leq (n - 1)C. \quad (20)$$

This bound strengthens the result of [CAE22], which considers the case where the trades for a specific CFMM can be reordered to maximize some payoff, by a factor of $\log(n)$. This graph is also known as the Cayley graph of the symmetric group generated by the transpositions, and, as one might expect, has deep connections to group theory (*cf.*, [Dia88]). We point out that the graph is $\binom{n}{2}$ -regular (as there are that many possible transpositions) and bipartite for future reference.

3.2 Smoothness over a graph

One simple way of thinking about the cost of MEV defined in (1) is as a measure of ‘smoothness’ of the function f over the set of permutations S , for a fixed set of transactions x . Given the graph $G = (V, E)$, we may define a reasonable notion of ‘smoothness’, as measured over the graph, by setting

$$C_G(f) = \sum_{\{i,j\} \in E} (f_i - f_j)^2.$$

The lower the value of $C_G(f)$ the more smooth a function f is over all ‘adjacent’ permutations. Note that, for any f this function is nonnegative since it is the sum of nonnegative terms. We will see that C_G and C are related, but C_G can sometimes provide a simpler, and more intuitive, description of ‘smoothness’ over permutations.

Properties. Like the original definition of C , the value of C_G is unchanged by offsets, *i.e.*,

$$C_G(f + \alpha \mathbf{1}) = C_G(f),$$

for $\alpha \in \mathbf{R}$, where $\mathbf{1} \in \mathbf{R}^{n_l}$ is the all-ones vector. Additionally, C_G is homogeneous of degree two in that

$$C_G(\alpha f) = \alpha^2 C_G(f),$$

for any $\alpha \in \mathbf{R}$.

Partial equivalence. The suggestive notation C_G here is in part justified by the following simple fact: for a fixed set of transactions $x \in A^n$, a function f has $C_G(f) = 0$ if, and only if, $C(f) = C(f, x) = 0$. (While the notation here is purposefully overloaded, the interpretations of each statement should be clear from context and the fact that x is fixed.) To see the forward implication, note that $C_G(f) = 0$ implies that

$$f_i = f_j, \quad \text{for every } \{i, j\} \in E. \quad (21)$$

Since the graph is connected, then $f_i = f_1$ for every $i = 1, \dots, n!$ and therefore f is a constant vector. Written another way, $f(\pi_i(x)) = f(x)$ for every permutation π_i , which exactly when the function is fair for transactions x , *i.e.*, when $C(f) = 0$. The reverse implication is immediate from the definitions. (We will show a tighter connection between C and C_G later in this section.)

3.2.1 The Laplacian of a graph

The function $C_G(f)$ can be written as the following homogeneous quadratic:

$$C_G(f) = f^T L f,$$

where the matrix $L \in \mathbf{R}^{n! \times n!}$ is defined as

$$L_{ij} = \begin{cases} -1 & \{i, j\} \in E \\ d_i & i = j, \end{cases}$$

and d_i is defined as the degree of node i ; *i.e.*, it is the number of vertices adjacent to node i . This matrix L is called the *Laplacian* and encodes a number of important properties of the graph G .

Eigenvectors and spectra. The matrix L is symmetric since $\{i, j\} = \{j, i\}$. A basic fact from linear algebra is that any symmetric matrix has an eigenvalue decomposition [Axl97]:

$$L = U \Sigma U^T,$$

where $\Sigma \in \mathbf{R}^{n! \times n!}$ is a diagonal matrix, while $U \in \mathbf{R}^{n! \times n!}$ is an orthogonal matrix with $U^T U = U U^T = I$. For the remainder of the paper, we will write $u_i \in \mathbf{R}^{n!}$ for the i th column of the matrix U and $\lambda_i = \Sigma_{ii}$. Since we may rearrange the columns of U and Σ as needed, we will also assume that $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_{n!}$. (This list of λ_i is called the *spectrum* of the graph G .) Using the definition of L and the u_i we have that

$$L u_i = \lambda_i u_i,$$

(which is why the above decomposition is called an eigenvalue decomposition) and therefore that

$$\lambda_i = u_i^T L u_i = C_G(u_i) \geq 0. \quad (22)$$

This implies that the λ_i are nonnegative. We can show that $\lambda_1 = 0$ since, using the definition of L ,

$$L\mathbf{1} = \mathbf{0}\mathbf{1} = 0,$$

where $\mathbf{1}$ is the all-ones vector, so we may assume that $u_1 = \mathbf{1}/\sqrt{n!}$. It is not hard to show that, since G is connected, $\lambda_2 > 0$.

Bounds on the largest eigenvalue. If the graph is a k -regular graph, such as the transposition graph with $k = \binom{n}{2}$, then we have the following bound on the largest eigenvalue:

$$\lambda_{n!} \leq 2k.$$

To see this, let \tilde{u} correspond to any eigenvector with eigenvalue $\tilde{\lambda} \geq 0$, and let $i \in \operatorname{argmax}_j |\tilde{u}_j|$ be its largest index, then:

$$\tilde{\lambda}|\tilde{u}_i| = |(L\tilde{u})_i| = \left| k\tilde{u}_i - \sum_{\{i,j\} \in E} \tilde{u}_j \right| \leq k|\tilde{u}_i| + \sum_{\{i,j\} \in E} |\tilde{u}_j| \leq k|\tilde{u}_i| + k|\tilde{u}_i| = 2k|\tilde{u}_i|.$$

The only major step is in the second inequality, which uses both the fact that the graph is k -regular (*i.e.*, every node has exactly k -neighbors) and the fact that $|\tilde{u}_i|$ is the largest entry in \tilde{u} . Since $\tilde{u} \neq 0$ then $\tilde{u}_i \neq 0$ by definition and canceling both sides we have

$$\tilde{\lambda} \leq 2k,$$

as required. (Using $\tilde{\lambda} = \lambda_{n!}$ and similarly for \tilde{u} gives the final result.) We will make use of this result, applied to the special case of the transposition graph, later in this section.

3.2.2 Fourier transform over graphs

We are now ready to define the *Fourier transform* of the function f over the graph G , given simply by the vector $\hat{f} \in \mathbf{R}^{n!}$, defined

$$\hat{f} = U^T f = \begin{bmatrix} u_1^T f \\ u_2^T f \\ \vdots \\ u_{n!}^T f \end{bmatrix},$$

where u_i is the i th eigenvector of the Laplacian L . Since $UU^T = I$ we can similarly write

$$f = U(U^T f) = U\hat{f} = \sum_{i=1}^{n!} \hat{f}_i u_i. \quad (23)$$

In this notation, it is clear that u_i act as a *Fourier basis* and each of the \hat{f}_i act as *Fourier coefficients* that weight the basis vectors. Indeed, from (22) and their definition, the u_i have

a reasonably simple interpretation: the u_i form an orthogonal basis with the property that, as i becomes larger, the u_i become ‘less smooth’ over G , or, in signal processing parlance, they represent ‘higher frequencies’ of the graph. We may then decompose f into its constituent ‘frequencies’, where \hat{f}_i gives the contribution of each frequency v_i to the overall function f . (Note that u_1 represents the lowest possible frequency: a constant.)

In this sense, the decomposition of (23) is very similar (and, in fact, in certain limits, for certain graphs, equivalent [Sin06]) to the traditional Fourier transform for signals. We have many of the usual identities, such as Parseval’s theorem, hold:

$$\|\hat{f}\|_2^2 = \hat{f}^T \hat{f} = \hat{f}^T (U^T U) \hat{f} = (U \hat{f})^T (U \hat{f}) = \|f\|_2^2, \quad (24)$$

since $U^T U = I$. We will also see next that the Fourier transform of f over G has deep connections with the definition of fairness presented in (1). As an important aside: the transform \hat{f} depends heavily on the structure of the graph G and may look very different for different graphs (even though the function f may be unchanged).

Largest eigenvector. For the case of the transposition graph, it is possible to explicitly construct its largest eigenvector. To do this, first note that the transposition graph is $\binom{n}{2}$ -regular. This means that its largest eigenvalue, $\lambda_{n!}$, is bounded by

$$\lambda_{n!} \leq 2 \binom{n}{2} = n(n-1),$$

from the previous discussion. Now, since, for any $v \in \mathbf{R}^{n!}$,

$$C_G(v) = v^T L v \leq \lambda_{n!} \|v\|_2^2,$$

we will exhibit a vector v with cost $C_G(v) = n(n-1) \|v\|_2^2$, which will show that $\lambda_{n!} = n(n-1)$ and that v is a corresponding eigenvector. We will define v in the following way

$$v_i = \begin{cases} +1 & \text{if } \pi_i \text{ is an even permutation} \\ -1 & \text{if } \pi_i \text{ is an odd permutation,} \end{cases}$$

where we have assumed that the i th node corresponds to permutation $\pi_i \in S$, for $i = 1, \dots, n!$. (For much more on the definition and existence of even/odd permutations, see [Con].) Now, note that

$$C_G(v) = \sum_{\{i,j\} \in E} (v_i - v_j)^2.$$

But, by definition of the transposition graph, two nodes i, j , are neighbors if, and only if, they differ by a transposition. So, if π_i is odd, then π_j must be even or vice versa, hence

$$C_G(v) = |E|(2)^2 = 4|E|,$$

where $|E|$ is the number of edges. Since the transposition graph is $\binom{n}{2}$ -regular, we know that

$$\frac{n! \binom{n}{2}}{2} = |E|,$$

and we also know that $\|v\|_2^2 = n!$, so, we get

$$C_G(v) = 4|E| = 2n! \binom{n}{2} = n(n-1)\|v\|_2^2,$$

as required. Normalizing this function gives an eigenvector corresponding to the largest eigenvalue,

$$u_n = \frac{1}{\sqrt{n!}}v.$$

(In fact, the largest eigenvalue has multiplicity one, so this eigenvector is unique, but we do not prove this here.)

Graph Fourier coherence. Having defined the Fourier transform of a graph, we now define a quantity that will be crucial in giving bounds on the fairness, following [PRSV18]. For every set of orthogonal eigenvectors u_j corresponding to a graph G , we define the *graph Fourier coherence* as:

$$\mu = \max_i \|u_i\|_\infty, \tag{25}$$

where $\|y\|_\infty = \max_i |y_i|$. This quantity measures how ‘concentrated’ the Fourier basis is when compared to the standard basis. Using basic estimates for the ℓ_∞ norm in the ℓ_2 norm since $\|u_i\|_2 = 1$:

$$\frac{1}{\sqrt{n!}} = \frac{1}{\sqrt{n!}} \|u_i\|_2 \leq \|u_i\|_\infty \leq \|u_i\|_2 = 1,$$

so $1/\sqrt{n!} \leq \mu \leq 1$. The upper bound occurs precisely when there is a vector u_i that is exactly aligned with one of the standard basis vectors, while the minimum value indicates that the Fourier basis is maximally ‘incoherent’ with the standard basis. As we will see, bounds on this quantity control our ability to localize a function f and its Fourier coefficients \hat{f} simultaneously. Intuitively, if μ is small, then the ‘size’ of f and \hat{f} must be traded off against each other in a sense that we will make precise. Note that, because μ is a graph-dependent quantity, the way that the edge set is constructed will affect this quantity, and hence our ability to bound the fairness. Note that the coherence (25) is a basis dependent phenomenon: if the eigenvalues are degenerate (*i.e.*, if there are repeated eigenvalues) then the coherence need not be unique and would depend on the specific eigenvectors chosen in the decomposition.

Coherence for connected graphs. In the case where a graph is connected and has more than one vertex, we can get a slightly better upper bound on the coherence of any eigenvector v , given by

$$\|v\|_\infty \leq \sqrt{1 - \frac{1}{n!}}.$$

In particular, the first eigenvector is known (the all-ones vector) and has

$$\frac{\|\mathbf{1}\|_\infty}{\sqrt{n!}} = \frac{1}{\sqrt{n!}},$$

which is exactly the lower bound on μ given above. Since the graph is connected, we must have that every other eigenvector v satisfies $\mathbf{1}^T v = 0$, from the previous discussion. For any nonzero vector, there must be at least one component that is strictly positive and one that is strictly negative; we write v^+ for the positive entries of v and v^- for the negative entries and assume, without loss of generality, that the maximum absolute value entry is in the positive part. (This is without loss of generality since, if v is an eigenvector, then $-v$ is as well.) With this, we have

$$\|v\|_\infty \leq \|v^-\|_1 \leq \|v^-\|_2 \sqrt{n! - 1} = \sqrt{1 - \|v^+\|_2^2} \sqrt{n! - 1} \leq \sqrt{1 - \|v\|_\infty^2} \sqrt{n! - 1}.$$

Rearranging and solving for $\|v\|_\infty$ gives the inequality above. The first inequality follows from the fact that the sums of the positive and negative entries must be equal, so the largest entry (which is positive) must be at least the sum of all of the negative entries. The second inequality follows from the fact that there are at most $n! - 1$ nonzero entries in v^- , while the third follows since $\|v\|_2 = 1$. Finally, the last follows from the fact that $\|v^+\|_\infty \leq \|v^+\|_2$ and that the largest entry is positive.

Note that this inequality is tight in that it can be saturated by choosing a vector with $v_1 = \sqrt{1 - 1/n!}$ and $v_i = -1/\sqrt{n!(n! - 1)}$ for $i = 2, \dots, n!$.

Coherence for the complete graph. We note that coherence is a phenomenon that depends on the specific Fourier basis that is chosen. (In particular, some graphs may have degenerate eigenvalues, in which case there are many possible eigendecompositions.) In the case of the complete graph, we have that there exists an eigendecomposition that saturates the above bound. This follows easily from the fact that any vector orthogonal to the all-ones vector is an eigenvector of the complete graph Laplacian. Choosing the vector above that saturates this inequality then gives us that, indeed, the coherence in any basis containing any such vector saturates the inequality for the complete graph.

Numerical estimates. We numerically compute μ for the various kinds of permutation graphs that are considered in §3.1. The results for complete graphs and transposition graphs are provided in table 1. Further numerical estimates for other kinds of graphs, including Erdős-Rényi graphs can be found in [PRSV18]. It can be seen that as the number of transpositions increases, the coherence decays the fastest for transposition graphs; this shows that the choice of the underlying graph can meaningfully affect the bounds on the cost of MEV, as we will see next. Code to reproduce these quantities is provided at

<https://github.com/bcc-research/spectrum-calculations>

Note that, in general, one might receive different numbers than those given in the table, as the transposition graph has degenerate eigenvalues. (The complete graph, on the other hand, is a global bound from the previous discussion.) The Fourier coefficients therefore also depend on the exact eigenvectors chosen which is why the graph coherence varies.

Number of transactions, n	Transposition graph	Complete graph
1	1.000	1.000
2	0.707	0.707
3	0.816	0.913
4	0.612	0.978
5	0.548	0.995
6	0.323	0.999
7	0.477	1.000

Table 1: Graph Fourier coherence, μ , for transposition and complete graphs.

3.3 Fairness and the Fourier transform

For the next section, we will assume that $\max_i f_i \geq \max_j -f_j$; in other words, the highest possible payoff is larger than the negative of the largest possible loss, which, using the translation invariance of C , may generally be assumed.

Cost of MEV. Given the above condition on f , note that $\max_i f_i = \|f\|_\infty$. Additionally, we may write the expectation of f uniformly over the possible permutations as $\mathbf{1}^T f / n!$. This means that we may rewrite the cost of MEV in (1) as (dropping the $x \in A^n$ as it is fixed):

$$C(f) = \|f\|_\infty - \frac{1}{n!} \mathbf{1}^T f.$$

Since $u_1 = \mathbf{1}/\sqrt{n!}$ and f can be decomposed as given in (23), then

$$\frac{1}{n!} \mathbf{1}^T f = \frac{1}{n!} \sum_{i=1}^{n!} \mathbf{1}^T u_i \hat{f}_i,$$

but, since $u_1^T u_i = \mathbf{1}^T u_i = 0$ unless $i = 1$ we then have

$$\frac{1}{n!} \mathbf{1}^T f = \frac{1}{\sqrt{n!}} \hat{f}_1,$$

so we can write

$$C(f) = \|f\|_\infty - \frac{1}{\sqrt{n!}} \hat{f}_1.$$

Lower bounds. This rewriting immediately provides a simple lower bound only in terms of \hat{f} , since $\sqrt{n!} \|f\|_\infty \geq \|f\|_2 = \|\hat{f}\|_2$, where the latter equality follows from (24), so

$$C(f) \geq \frac{1}{\sqrt{n!}} \left(\|\hat{f}\|_2 - \hat{f}_1 \right). \quad (26)$$

We can view the term $\|\hat{f}\|_2 - \hat{f}_1$ roughly as asking ‘how much more do the higher frequencies of f contribute over \hat{f}_1 ’? Alternatively: how much of the mass of \hat{f} is contained in all of the nonconstant frequencies? This also gives an interesting observation. If f is a perfectly fair function, we know that $C(f) = 0$. Using (26), we get $\|\hat{f}\|_2 \leq \hat{f}_1$, or, after squaring both sides and subtracting,

$$\hat{f}_2^2 + \hat{f}_3^2 + \cdots + \hat{f}_{n!}^2 \leq 0,$$

which immediately implies that \hat{f}_1 is the only possible nonzero coefficient. This corresponds neatly to the ‘usual’ notion that a constant function has a Fourier transform whose only support is at the lowest possible frequency. (We show the ‘opposite’, that worst-cost functions have nonzero coefficients at high frequencies, later.)

Upper bounds. A basic upper bound for this problem follows from the fact that

$$\|f\|_\infty = \left\| \sum_{i=1}^{n!} \hat{f}_i u_i \right\|_\infty \leq \sum_{i=1}^{n!} |\hat{f}_i| \|u_i\|_\infty \leq \|\hat{f}\|_1 \left(\max_i \|u_i\|_\infty \right) = \|\hat{f}\|_1 \mu,$$

which gives

$$C(f) \leq \mu \|\hat{f}\|_1 - \frac{1}{\sqrt{n!}} \hat{f}_1.$$

Worst-cost functions on the transposition graph. In a very specific sense, the ‘worst-cost’ functions defined in the previous sections are some of the ‘spikiest’ functions over the graph. For example, in the special case of the transposition graph, the largest eigenvector is always nonzero at every node of the graph, which immediately shows that

$$\hat{e}_i = e_i^T u_{n!} = \pm 1/\sqrt{n!},$$

so these worst-cost functions always have nonzero components at the highest possible frequency. Much in the same way that the (usual) Fourier transform of very ‘spiky’ functions has high frequency components, we see that the graph Fourier transform of a worst-cost function (namely, the basis vectors) similarly has high frequency components when considering the transposition graph.

3.3.1 Relationship between C and C_G

Unsurprisingly, there is a relationship between the cost of fairness C and the measure of ‘smoothness’ C_G , given by,

$$\frac{1}{\sqrt{\lambda_{n!} n!}} \sqrt{C_G(f)} \leq C(f) \leq \sqrt{\frac{C_G(f)}{\lambda_2}}.$$

We show the two inequalities that relate these two functions in what follows.

Lower bound. There is a lower bound in terms of the largest eigenvalue given by

$$C(f) \geq \frac{1}{\sqrt{\lambda_{n!}n!}} \sqrt{C_G(f)}.$$

To see that the original bound is true, note that, since $\lambda_{n!}$ is the largest eigenvalue of L , we have

$$C_G(f) = f^T L f \leq \lambda_{n!} \|f\|_2^2.$$

Using (26) and the fact that $\|f\|_2 = \|\hat{f}\|_2$ we then have

$$C(f) \geq \frac{1}{\sqrt{n!}} \left(\sqrt{\frac{C_G(f)}{\lambda_{n!}}} - \hat{f}_1 \right).$$

Finally, using the fact that both C and C_G are translation invariant, set $\alpha = \hat{f}_1/\sqrt{n!}$ and note that the first Fourier coefficient of the function $f - \alpha \mathbf{1}$ is zero to get

$$C(f) = C(f - \alpha \mathbf{1}) \geq \frac{1}{\sqrt{n!}} \sqrt{\frac{C_G(f - \alpha \mathbf{1})}{\lambda_{n!}}} = \frac{1}{\sqrt{\lambda_{n!}n!}} \sqrt{C_G(f)}.$$

Upper bound. We also have an upper bound when $\mathbf{1}^T f \geq 0$ given by

$$C(f) \leq \sqrt{\frac{C_G(f)}{\lambda_2}}.$$

This is a refinement of the previous claim that $C_G(f) = 0$ implies that $C(f) = 0$, since we know that $\lambda_2 > 0$ as the graph is connected. (Note that this inequality fails when the graph is not connected as we would have $\lambda_2 = 0$.) To see this, bound $C(f)$ by

$$C(f) = \max_i |f_i| - \mathbf{1}^T f / n! \leq \max_i |f_i - \mathbf{1}^T f / n!| \leq \|f - \alpha \mathbf{1}\|_2,$$

where we have defined $\alpha = \mathbf{1}^T f / n!$. We will overload notation slightly and write \hat{f} for the Fourier coefficients of $f - \alpha \mathbf{1}$. The first coefficient of this function, \hat{f}_1 , is zero since it is orthogonal to the all-ones vector by definition of α , $\mathbf{1}^T (f - \alpha \mathbf{1}) = 0$. This means that

$$C_G(f - \alpha \mathbf{1}) = \sum_{i=2}^{n!} \lambda_i \hat{f}_i^2 \geq \lambda_2 \sum_{i=2}^{n!} \hat{f}_i^2 = \lambda_2 \|\hat{f}\|_2^2,$$

as the eigenvalues are in nondecreasing order, by assumption, and $\hat{f}_1 = 0$. Finally, using the fact that $\|\hat{f}\|_2 = \|f - \alpha \mathbf{1}\|_2$ and our previous inequality for C , we get

$$C(f) \leq \sqrt{\frac{C_G(f - \alpha \mathbf{1})}{\lambda_2}} = \sqrt{\frac{C_G(f)}{\lambda_2}},$$

where the last equality follows from the fact that C_G is translation-invariant. There is a (relatively weak) lower bound for λ_2 due to [Moh91], which is

$$\lambda_2 \geq \frac{4}{\mathbf{diam}(G)n!}.$$

This would imply the following bound, which does not depend on the eigenvalues of the Laplacian:

$$C(f) \leq \frac{\sqrt{\mathbf{diam}(G)n!}}{2} \sqrt{C_G(f)}.$$

Transposition graph. In the special case of the transposition graph, we know that it is $\binom{n}{2}$ -regular and has diameter $n - 1$, which gives the following relationship between C and C_G ,

$$\frac{1}{\sqrt{n(n-1)n!}} \sqrt{C_G(f)} \leq C(f) \leq \frac{\sqrt{(n-1)n!}}{2} \sqrt{C_G(f)}.$$

Note that these bounds are particularly loose and differ by large constants. Since the transposition graph has a lot of structure that we do not exploit here, it is very likely that the bounds provided can be significantly improved.

4 Conclusion

We constructed a simple framework that may be used to assess the economic impact of reordering MEV, which is excess value that a monopolistic validator can extract from a decentralized network. We define the *cost of MEV*, which is the difference between the worst-case and average case payoff to the validator. Our definition allows us to compute a notion of ‘fairness’ for a set of transaction orderings that is a function of the economic value that can be extracted from users. This is in contrast to definitions of ‘fair ordering’ that define ‘fairness’ according to arrival times of transactions.

The key objects studied are payoff functions, which are real-valued functions that map transactions and their order to a utility, or payoff, realized by the validator. We demonstrate that all payoff functions realized by users can be written as linear combinations of the worst-cost payoff functions, which are similar in spirit to liquidations—core mechanisms within decentralized finance.

We then make the distinction between ‘spiky’ and ‘flat’ payoff functions more quantitative in two different directions: via simple bounds and via the Fourier transform on a graph. We show that for functions that are smooth, so too is the difference between their worst and average case behavior. Additionally, we noted that payoff functions on permutations may also be seen analogously as functions defined on graphs, where the vertices are the permutations. In this case, we showed that the spectra of these functions provide non-trivial upper and lower bounds on the cost of MEV, analyzed via the Fourier transform on graphs.

Future work. There are various directions of future work. First, the Fourier bounds provided in §3 can be sharpened by utilizing the representation theory of the symmetric group. Representation theory allows for decomposing the graph Laplacian into the direct sum of smaller matrices, each of which can have sharper bounds than those found here. Prior work on analyzing the fairness of ranked voting [Dia88, Mos22] can likely be modified and used to directly improve the bounds of §3.

Another direction of work that one can take is to study ‘hierarchical’ costs of MEV. A number of new MEV auction designs such as SUAVE [Fla22], Jito [Lab23], Anoma [SY22], and Skip [HPM23] have auctions on a per domain or application level. A domain corresponds to a particular set of contracts or code where reordering is allowed, while reordering between domains is not controlled by any individual party (such as a monopolist sequencer or validator) [OSS⁺21]. In such a world, one can imagine the set of n pending transactions as distributed into k groups of transactions, representing the k distinct domains. One can compute a cost of MEV for each group, which leads to a natural question of what the aggregate cost of MEV is over all domains. This open problem may have some analogues to the current framework, which we leave for future work.

Acknowledgments

The authors would like to thank Peteris Erins for suggestions, comments, and edits on earlier drafts of this paper.

References

- [AAE⁺22] Guillermo Angeris, Akshay Agrawal, Alex Evans, Tarun Chitra, and Stephen Boyd. Constant Function Market Makers: Multi-asset Trades via Convex Optimization. In *Handbook on Blockchain*, page 31. Springer, first edition, 2022.
- [Axl97] Sheldon Axler. *Linear algebra done right*. Springer Science & Business Media, 1997.
- [BGR23] Maryam Bahrani, Pranav Garimidi, and Tim Roughgarden. Transaction fee mechanism design with active block producers. *arXiv preprint arXiv:2307.01686*, 2023.
- [BO22] Joseph Bebel and Dev Ojha. Ferveo: Threshold decryption for mempool privacy in bft networks. *Cryptology ePrint Archive*, 2022.
- [BZ23] Massimo Bartoletti and Roberto Zunino. A theoretical basis for blockchain extractable value. *arXiv preprint arXiv:2302.02154*, 2023.
- [CAE22] Tarun Chitra, Guillermo Angeris, and Alex Evans. Differential Privacy in Constant Function Market Makers. *Financial Cryptography and Data Security 2022*, 2022.

- [CK22] Tarun Chitra and Kshitij Kulkarni. Improving proof of stake economic security via mev redistribution. In *Proceedings of the 2022 ACM CCS Workshop on Decentralized Finance and Security*, pages 1–7, 2022.
- [Cla84] Allan Clark. *Elements of abstract algebra*. Courier Corporation, 1984.
- [Con] Keith Conrad. The sign of a permutation. *Online*.
- [DGK⁺20] Philip Daian, Steven Goldfeder, Tyler Kell, Yunqi Li, Xueyuan Zhao, Iddo Bentov, Lorenz Breidenbach, and Ari Juels. Flash boys 2.0: Frontrunning in decentralized exchanges, miner extractable value, and consensus instability. In *2020 IEEE Symposium on Security and Privacy (SP)*, pages 910–927. IEEE, 2020.
- [Dia88] Persi Diaconis. Group representations in probability and statistics. *Lecture notes-monograph series*, 11:i–192, 1988.
- [Fla22] Flashbots. The future of mev is suave: Flashbots, Nov 2022.
- [Fla23] Flashbots Team. Mev explore, 2023.
- [GSBV] Jacob Greene, Ashar Shahid, Burak Benligiray, and Heikki Vanttinen. Oracle extractable value (oev) through order flow auctions.
- [HPM23] Sam Hart, Barry Plunkett, and Magnus Mareneck. x/builder: The first Sovereign MEV module for protocol-owned building — ideas.skip.money. <https://ideas.skip.money/t/x-builder-the-first-sovereign-mev-module-for-protocol-owned-building/57>, 2023.
- [JDE⁺23] Nicholas AG Johnson, Theo Diamandis, Alex Evans, Henry de Valence, and Guillermo Angeris. Concave pro-rata games. *arXiv preprint arXiv:2302.02126*, 2023.
- [KDC22] Kshitij Kulkarni, Theo Diamandis, and Tarun Chitra. Towards a theory of maximal extractable value i: Constant function market makers. *arXiv preprint arXiv:2207.11835*, 2022.
- [KDL⁺21] Mahimna Kelkar, Soubhik Deb, Sishan Long, Ari Juels, and Sreeram Kannan. Themis: Fast, strong order-fairness in byzantine consensus. *Cryptology ePrint Archive*, 2021.
- [Lab23] Jito Labs. Jito Block Engine Expands Access to All Solana MEV Traders Jito Labs — jito.wtf. <https://www.jito.wtf/blog/jito-block-engine-expands-access-to-all-solana-mev-traders/>, 2023.
- [Mac23] Andrew W Macpherson. Adversarial queues and trading on a cfmm. *arXiv preprint arXiv:2302.01663*, 2023.

- [MDFO22] Conor McMenamin, Vanesa Daza, Matthias Fitzi, and Padraic O’Donoghue. Fairtradex: A decentralised exchange preventing value extraction. In *Proceedings of the 2022 ACM CCS Workshop on Decentralized Finance and Security*, pages 39–46, 2022.
- [Moh91] Bojan Mohar. Eigenvalues, diameter, and mean distance in graphs. *Graphs and combinatorics*, 7(1):53–64, 1991.
- [Mos22] Elchanan Mossel. Probabilistic view of voting, paradoxes, and manipulation. *Bulletin of the American Mathematical Society*, 59(3):297–330, 2022.
- [MS22] Dahlia Malkhi and Pawel Szalachowski. Maximal extractable value (mev) protection on a dag. *arXiv preprint arXiv:2208.00940*, 2022.
- [OSS⁺21] Alexandre Obadia, Alejo Salles, Lakshman Sankar, Tarun Chitra, Vaibhav Chelani, and Philip Daian. Unity is strength: A formalization of cross-domain maximal extractable value. *arXiv preprint arXiv:2112.01472*, 2021.
- [PRSV18] Nathanael Perraudin, Benjamin Ricaud, David I Shuman, and Pierre Vandergheynst. Global and local uncertainty principles for signals on graphs. *AP-SIPA Transactions on Signal and Information Processing*, 7, 2018.
- [QZG22] Kaihua Qin, Liyi Zhou, and Arthur Gervais. Quantifying blockchain extractable value: How dark is the forest? In *2022 IEEE Symposium on Security and Privacy (SP)*, pages 198–214. IEEE, 2022.
- [Res23] Max Resnick. Contingent fees in order flow auctions. *arXiv preprint arXiv:2304.04981*, 2023.
- [RK23] Antoine Rondelet and Quintus Kilbourn. Threshold encrypted mempools: Limitations and considerations. *arXiv preprint arXiv:2307.10878*, 2023.
- [Rou16] Tim Roughgarden. *Twenty Lectures on Algorithmic Game Theory*. Cambridge University Press, 2016.
- [Sin06] A. Singer. From graph to manifold laplacian: The convergence rate. *Applied and Computational Harmonic Analysis*, 21(1):128–134, 2006. Special Issue: Diffusion Maps and Wavelets.
- [SY22] Awa Sun Yin. An overview of anoma’s architecture. Anoma, 2022.
- [XFP23] Matheus Venturyne Xavier Ferreira and David C Parkes. Credible decentralized exchange design via verifiable sequencing rules. In *Proceedings of the 55th Annual ACM Symposium on Theory of Computing*, pages 723–736, 2023.
- [ZQT⁺21] Liyi Zhou, Kaihua Qin, Christof Ferreira Torres, Duc V Le, and Arthur Gervais. High-frequency trading on decentralized on-chain exchanges. In *2021 IEEE Symposium on Security and Privacy (SP)*, pages 428–445. IEEE, 2021.